

# Efficient and Distributed training with TensorFlow on Piz Daint

Synchronous Distributed Training with TensorFlow and Horovod

Rafael Sarmiento and Guilherme Peretti-Pezzi


ETHZürich / CSCS

Lugano, 14<sup>th</sup>-15<sup>th</sup> March 2019

# Outline

- Stochastic Gradient Descent
- [lab] Simple Stochastic Gradient Descent
- Synchronous Distributed Stochastic Gradient Descent
- Ring Allreduce
- Horovod
- [lab] Simple Stochastic Gradient Descent with Horovod
- [lab] CNN models with `tf.keras` + Horovod
- [lab] CNN models with TensorFlow's Estimator + Horovod

# We want to train a model on this data



# We choose a model and a cost function

$$y = mx + n$$

$$L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$$


# We choose a model and a cost function

$$y = mx + n$$

$$L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$$

$$L = \frac{1}{N} \sum_i^N (mx_i + n - y_i)^2$$



# We choose a model and a cost function





$$y = mx + n$$

$$L = \frac{1}{N} \sum_i^N (mx_i + n - y_i)^2$$

# We need to choose an optimizer




# We need to choose an optimizer






# <Stochastic> Gradient Descent




- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  for a batch of  $N$  samples  $\{x, y\}$  (forward pass)

# <Stochastic> Gradient Descent




- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  for a batch of  $N$  samples  $\{x, y\}$  (forward pass)
- Compute the gradients of the loss function with respect to the parameters of the model  $\frac{\partial L}{\partial W} \big|_{\{x, y\}}$  (backpropagation)

# <Stochastic> Gradient Descent



- Evaluate the loss function  $L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$  for a batch of  $N$  samples  $\{x, y\}$  (forward pass)
- Compute the gradients of the loss function with respect to the parameters of the model  $\frac{\partial L}{\partial W} \big|_{\{x, y\}}$  (backpropagation)
- Update the parameters  $W_t = W_{t-1} - \eta \frac{\partial L}{\partial W} \big|_{\{x, y\}_{t-1}}$


# <Stochastic> Gradient Descent



Gradient  
Descent


```
batch_size = training_set_size
```

# <Stochastic> Gradient Descent



Gradient  
Descent


`batch_size = training_set_size`



Stochastic Gradient  
Descent


`batch_size = 1`

# <Stochastic> Gradient Descent




Gradient  
Descent

`batch_size = training_set_size`



Stochastic Gradient  
Descent

`batch_size = 1`



Minibatch Stochastic Gradient  
Descent


`1 < batch_size < training_set_size`

# [lab] Simple Stochastic Gradient Descent

Let's run the notebook `SGD/simple_SGD_with_custom_model.ipynb`. There we use an unidimensional linear model to understand the trajectories of the SGD minimization.


Try different batch sizes and see how the trajectory changes.

# <Stochastic> Gradient Descent




Gradient  
Descent

`batch_size = training_set_size`



Stochastic Gradient  
Descent

`batch_size = 1`



Minibatch Stochastic Gradient  
Descent

`1 < batch_size < training_set_size`



- The batch size is a hyperparameter


- The batch size is a hyperparameter
- Large batches may not fit on the GPU memory

- The batch size is a hyperparameter
- Large batches may not fit on the GPU memory
- Splitting the training into multiple workers enables the use of large batches

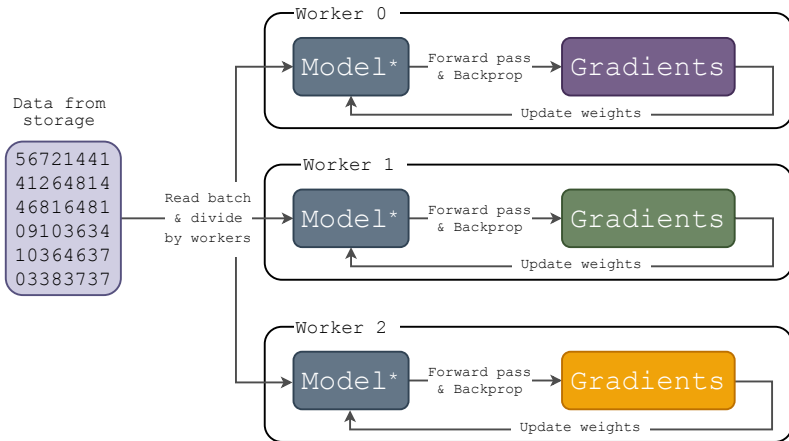
- The batch size is a hyperparameter
- Large batches may not fit on the GPU memory
- Splitting the training into multiple workers enables the use of large batches
- A large batch size does not necessarily mean faster convergence

# Distributing the training with data parallelism


Data from  
storage



# Distributing the training with data parallelism



# Distributing the training with data parallelism




# The Allreduce operation

- The Allreduce name comes from the MPI standard.
- MPI defines the function `MPI_Allreduce` to reduce values from all ranks and broadcast the result of the reduction such that all processes have a copy of it at the end of the operation.
- Allreduce can be implemented in different ways depending on the problem.



# Ring Allreduce




---

<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow


# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow

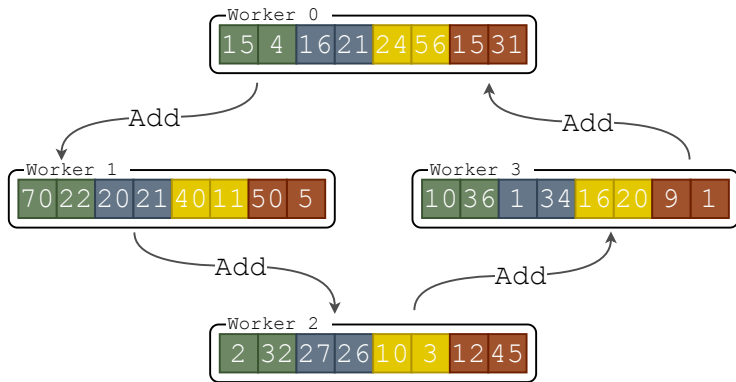
# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow

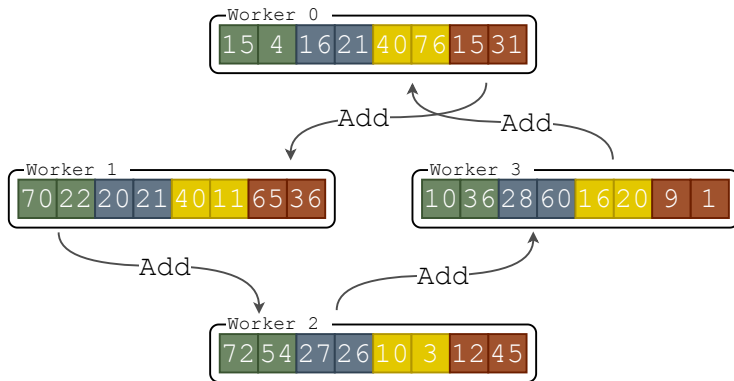
# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow


# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow


# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow

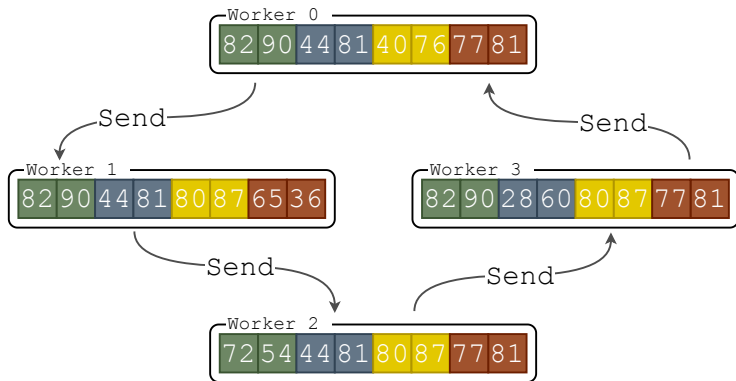
# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow

# Ring Allreduce

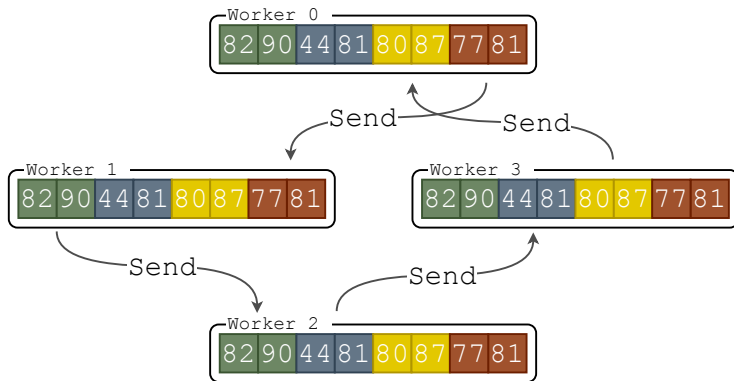


<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow



# Ring Allreduce



<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow

# Ring Allreduce

- Each of the  $N$  workers communicates only with other two workers  $2(N - 1)$  times.
- The values of the reduction are obtained with the first  $N - 1$  communications.
- The second  $N - 1$  communications are performed to update the reduced values on all workers.
- The total amount of data sent by each worker  $\left[ 2(N - 1) \frac{\text{ArraySize}}{N} \right]$  is virtually independent of the number of workers .

---

<sup>1</sup> [Baidu-Allreduce on GitHub](#)

<sup>2</sup> [A. Sergeev, M. del Balse. Horovod: fast and easy distributed deep learning in TensorFlow](#)

# Communication between Cray XC50 Nodes on Piz Daint

- Aries interconnect with the Dragonfly topology
- Direct communications between nodes on the same electrical group (2 cabinets / 384 nodes)
- Communications between nodes on different electrical groups passes by switches (submit with option `#SBATCH --switches=1` to make your job wait for a single-group allocation)
- More info on [CSCS user portal](#)

# Horovod



*Horovod is an open-source distributed training framework for TensorFlow, Keras, PyTorch, and MXNet developed by Uber. The goal of Horovod is to make distributed Deep Learning fast and easy to use.*

# Horovod



- Minimal code modification required
- Uses bandwidth-optimal communication protocols
- Seamless integration with Cray-MPICH and use the NVidia Collective Communications Library (NCCL-2)
- Actively developed
- Growing community

# NVIDIA Collective Communications Library (NCCL)



*NCCL implements multi-GPU and multi-node collective communication primitives that are performance optimized for NVIDIA GPUs. NCCL provides routines such as Allgather, Allreduce and Broadcast, optimized to achieve high bandwidth over PCIe and NVLink high-speed interconnect.*

# Running TensorFlow + Horovod on Piz Daint

```
#!/bin/bash -l
#SBATCH --job-name=tf_hvd
#SBATCH --time=00:15:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-core=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=12
#SBATCH --constraint=gpu

module load daint-gpu
module load Horovod/0.16.0-CrayGNU-18.08-tf-1.12.0
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK
export NCCL_DEBUG=INFO
export NCCL_IB_HCA=ipogif0
export NCCL_IB_CUDA_SUPPORT=1

srun python my_script.py
```

# Horovod: 1. Initialize the library (TensorFlow)

```
import horovod.tensorflow as hvd  
hvd.init()
```



# Horovod: 1. Initialize the library (tf.keras)

```
import horovod.tensorflow.keras as hvd  
hvd.init()
```


## Horovod: 2. Sync initial state among workers




## Horovod: 2. Sync initial state among workers




## Horovod: 2. Sync initial state among workers




## Horovod: 2. Sync initial state among workers



## Horovod: 2. Sync initial state among workers



## Horovod: 2. Sync initial state among workers



## Horovod: 2. Sync initial state among workers (TensorFlow)

```
hooks = [hvd.BroadcastGlobalVariablesHook(0)]  
with tf.train.MonitoredTrainingSession(hooks=hooks, ...) as sess:  
    ...
```



## Horovod: 2. Sync initial state among workers (TensorFlow - Estimator API)

```
estimator = tf.estimator.Estimator(...)

hooks = [hvd.BroadcastGlobalVariablesHook(0)]
estimator.train(input_fn=train_input_fn,
                 steps=NUM_STEPS,
                 hooks=hooks)
```

## Horovod: 2. Sync initial state among workers (tf.keras)

```
callbacks = [hvd.callbacks.BroadcastGlobalVariablesCallback(0)]  
model.fit(dataset, ..., callbacks=callbacks)
```

# Horovod: 3. Checkpoints

```
# Save checkpoints for the worker of rank 0.  
# This will prevent all workers from corrupting a  
# single checkpoint file.  
if hvd.rank() == 0:  
    ...
```




## Horovod: 4. Wrap optimizer with Horovod's distributed one (TensorFlow)

```
opt = tf.train.GradientDescentOptimizer(learning_rate)
opt = hvd.DistributedOptimizer(opt)
```

## Horovod: 4. Wrap optimizer with Horovod's distributed one (`tf.keras`)

```
opt = tf.keras.optimizers.SGD(learning_rate)
opt = hvd.DistributedOptimizer(opt)
```

# Benchmarks results on Piz Daint (CNNs on Imagenet)



# Some additional considerations

- Data must be split equally by workers to avoid load imbalance.
- If applicable, data can be split such that each worker does not need to read all files.
- Dataset splits resulting in non-homogeneous datasets may harm the convergence.
- Consider scaling the learning rate (`learning_rate * hvd.size()`)

# Intuition on scaling the learning rate

$$L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$$

$$W_{t+1} = W_t - \eta \frac{\partial L}{\partial W} \Big|_{\{x,y\}_t}$$

---

<sup>1</sup> P. Goyal et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour



# Intuition on scaling the learning rate

$$L = \frac{1}{N} \sum_i^N l(\hat{y}_i, y_i)$$

$$W_{t+1} = W_t - \eta \frac{\partial L}{\partial W} \Big|_{\{x,y\}_t}$$

$$W_{t+1} = W_t - \frac{\eta}{N} \sum_{i \in t}^N \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

---

<sup>1</sup> P. Goyal et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour

# Intuition on scaling the learning rate

$$W_{t+k} = W_t - \frac{\eta}{N} \sum_j^k \sum_{i \in t_j}^N \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

---

<sup>1</sup> P. Goyal et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour

# Intuition on scaling the learning rate

$$W_{t+k} = W_t - \frac{\eta}{N} \sum_j^k \sum_{i \in t_j}^N \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

$$W_{t+1} = W_t - \frac{\eta}{kN} \sum_{i \in t}^{kN} \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

---

<sup>1</sup> P. Goyal et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour

# Intuition on scaling the learning rate

$$W_{t+k} = W_t - \frac{\eta}{N} \sum_j^k \sum_{i \in t_j}^N \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

$$W_{t+1} = W_t - \frac{\eta}{kN} \sum_{i \in t}^{kN} \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

$$W_{t+1} = W_t - \frac{k\eta}{kN} \sum_{i \in t}^{kN} \frac{\partial l}{\partial W} \Big|_{\{x,y\}_i}$$

---

<sup>1</sup> P. Goyal et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour

# TensorFlow's distribution strategy

- TensorFlow-1.12.0 includes support for synchronous distributed training using Ring Allreduce with `collectiveAllReduceStrategy` (already available, but still under development)
- Currently it involves the definition of an environment variable which is different for each worker

```
TF_CONFIG='{  
    "cluster": {"worker": ["IP_NODE1:PORT", "IP_NODE2:PORT"]},  
    "task": {"type": "worker", "index": WORKER_RANK}  
}'
```

- It looks forward to involve as little code modification as possible
- An example is included with the lab material

# [lab] Simple Stochastic Gradient Descent with Horovod

The script `hvd_simple_SGD_with_custom_model_exercise.py` uses the same model that we saw before on the notebook. We will adapt it to Horovod and we will run it with 2 workers.

```
# login with the -X option to open plot windows
ssh -X studxx@ela.cscs.ch
ssh -X studxx@daint.cscs.ch

alloc -C gpu -N 2 --res tensor2
module load daint-gpu
module load Horovod/0.16.0-CrayGNU-18.08-tf-1.12.0
cd tensorflow-training-cscs/SGD
srun python hvd_simple_SGD_with_custom_model_exercise.py
```

The script will save the trajectories of the two workers that can be visualized simply by running

```
python plot_hvd_simple_SGD_with_custom_model.py
```

Visualize the trajectories before and after adding each Horovod modification.

# [lab] CNN models with `tf.keras` + Horovod

On `models_from_keras_applications` there are scripts to do simple training of popular Convolutional Neural Networks models on ImageNet. The scripts starting with `hvd_` contain Horovod code, while the other ones contain the equivalent single-node code. Adapt them to Horovod and run them in two nodes.

```
# login with SSH
ssh studxx@ela.cscs.ch
ssh studxx@daint.cscs.ch

cd models_from_keras_applications

alloc -C gpu -N 2 --res tensor2
module load daint-gpu
module load Horovod/0.16.0-CrayGNU-18.08-tf-1.12.0
srun python keras_<model>_imagenet.py
```

# [lab] CNN models with TensorFlow's Estimator + Horovod

On `custom_estimators_from_benchmaks_models` there are scripts to do simple training of popular Convolutional Neural Networks models on ImageNet. The scripts starting with `hvd_` contain Horovod code, while the other ones contain the equivalent single-node code. Addapt them to Horovod and run them in two nodes.

```
cd custom_estimators_from_benchmaks_models
# download the models from https://github.com/tensorflow/benchmarks
bash get_models_from_tfbenchmarks.sh
cd models_from_benchmark

alloc -C gpu -N 2 --res tensor2
module load daint-gpu
module load Horovod/0.16.0-CrayGNU-18.08-tf-1.12.0
cp ../estimator_<model>_imagenet.py .
srun python estimator_<model>_imagenet.py
```



Thank you for your attention!