

Efficient and Distributed training with TensorFlow on Piz Daint

Input pipelines with TensorFlow's `tf.data` API

Rafael Sarmiento and Guilherme Peretti-Pezzi


ETHZürich / CSCS

Lugano, 14th-15th March 2019

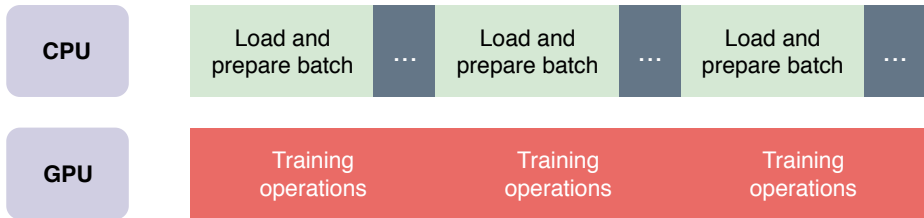
Outline

- Input pipelines
- [lab] Building input pipelines
- Read and write TFRecord files
- [lab] Reading and writing TFRecord files
- [lab] Decoding ImageNet data from a TFRecord file
- Optimizing pipelining
- Feeding datasets to models
- [lab] Feeding datasets to models

Consider large shared datasets



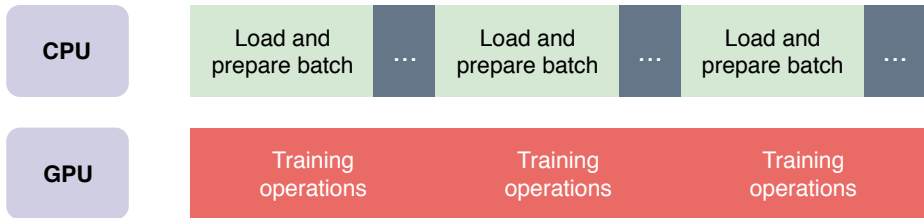
Pipelining



Pipelining



Pipelining




Input pipelines



Extract

Read data from persistent storage (HDD, SSD, GCS, HDFS, ...).



Transform

Use **CPU cores** to parse and perform preprocessing operations on the data, shuffling and batching.



Load

Load the transformed data onto the accelerator devices that execute the model.

TensorFlow's `tf.data` API

Extract

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

Transform

```
dataset = dataset.shuffle(1000)
```

```
dataset = dataset.batch(64)
```

```
dataset = dataset.repeat(100)
```

Load

```
iterator = dataset.make_one_shot_iterator()
```

```
next_item = iterator.get_next()
```


TensorFlow's `tf.data` API: Extract

```
# Read data from TensorFlow's TFRecord file
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

TensorFlow's `tf.data` API: Extract

```
# Read data from csv file:
#
# column1      column2      column3
#   0.0         4.1         dog
#   3.2         5.7         cat
dataset = tf.data.experimental.CsvDataset('train.csv',
                                           header=True,
                                           record_defaults=[tf.float32,
                                                             tf.float32,
                                                             tf.string])
```

TensorFlow's `tf.data` API: Extract

```
# Read data from numpy arrays or `tf.Tensor`s in memory:  
dataset = tf.data.Dataset.from_tensor_slices((x_numpy, y_numpy))
```

TensorFlow's `tf.data` API: Extract

```
# Read data with custom logic defined as a generator:
# def dataset_generator():
#     f = open('train.txt')
#     lines = f.readlines()
#     for l in lines:
#         x = np.array(l.split()[:4]).astype(np.float32)
#         y = np.array(l.split()[4]).astype(np.int32)
#         yield x, y
dataset = tf.data.Dataset.from_generator(dataset_generator,
                                         output_types=(tf.float32,
                                                         tf.int32))
```

TensorFlow's `tf.data` API: Transform

```
# Reading data
```

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```



TensorFlow's `tf.data` API: Transform

```
# Shuffle groups of 4 components at the time  
dataset = dataset.shuffle(4)
```

C D A B H E F G J L K I O M N P

TensorFlow's `tf.data` API: Transform

```
# Group data in batches of 4 components  
dataset = dataset.batch(4)
```



TensorFlow's `tf.data` API: Transform

```
# Filter out elements for which the user-defined function  
# `filter_vowels` returns False.  
dataset = dataset.filter(filter_vowels)
```




TensorFlow's `tf.data` API: Transform

```
# Operations like map or filter can be applied batch by batch  
# to take advantage of vectorial operations.  
# Here we apply the user-defined function `change_color` to  
# each of the components of the dataset batch by batch.  
dataset = dataset.map(change_color)
```



TensorFlow's `tf.data` API: Transform

```
# Repeat twice the dataset to make 2 epochs  
dataset = dataset.repeat(2)
```



TensorFlow's `tf.data` API: Transform

```
# Reading data
```

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

A B C D E F G H I J K L M N O P

```
# Split the input pipeline by shards
```

```
dataset = dataset.shard(num_ranks, rank)
```

A C E G I K M O (Rank 0)

B D F H J L N P (Rank 1)

TensorFlow's `tf.data` API: Transform

```
# listing input files
```

```
dataset = tf.data.Dataset.list_files(['file_1', 'file_2'], <opts>)
```



Content of file_1

Content of file_2

```
# interleaving with block length 1
```

```
dataset = dataset.interleave(dataset_reader, block_length=1, <opts>)
```



TensorFlow's `tf.data` API: Transform

```
# listing input files
```

```
dataset = tf.data.Dataset.list_files(['file_1', 'file_2'], <opts>)
```



Content of file_1

Content of file_2

```
# interleaving with block length 2
```

```
dataset = dataset.interleave(dataset_reader, block_length=2, <opts>)
```



TensorFlow's `tf.data` API: Transform

... and much more!

TensorFlow's `tf.data` API: Load

```
# Create a stream of data from the hard drive to the model  
iterator = dataset.make_one_shot_iterator()  
next_item = iterator.get_next()
```

TensorFlow's `tf.data` API: Load

```
# Create a stream of data from the hard drive to the model  
iterator = dataset.make_one_shot_iterator()  
features, label = iterator.get_next()
```


Iterating over data

```
with tf.Session() as sess:
    try:
        while True:
            features, label = sess.run(next_item)
            print('features: %s | label: %s' % (features, label))
    except tf.errors.OutOfRangeError:
        print('The dataset ran out of entries!')
```



```
# features: [[0.32762216 0.19132466]] | label: [2]
# features: [[0.40871843 0.02722579]] | label: [1]
# ...
# features: [[0.13172416 0.40897961]] | label: [1]
# The dataset ran out of entries!
```

[lab] Reading and writing TFRecord files

Let's open an empty notebook and create some simple input pipelines. A good starting point can be to generate random numpy data and create the pipeline including maps, filters, batching, shuffling and repeating operations. `tf.data.Dataset.from_tensor_slices` can be used for the *extract* phase.

Then we can run together the notebooks `getting_started_with_tensorflows_dataset_api_*.ipynb`, that are on the folder `input_pipelines/` to try other examples.

TFRecord file system

- TFRecord is a simple record-oriented binary format
- Data is stored as collections of meaningful units (records) in contrast to a byte-oriented filesystem, where the data is treated as an unformatted stream of bytes
- On the deep learning context each record would be an item of the dataset
- TFRecord is the format of data storage recommended for TensorFlow

TFRecord file system

```
dataset = tf.data.TFRecordDataset("./train.tfrecords")
```

Write data as a TFRecord

```
# For representing data structures, for instance images.  
def _bytes_feature(value):  
    return tf.train.Feature(bytes_list=tf.train.BytesList(value=[value]))  
  
# For representing integer values, for instance integer labels.  
def _int64_feature(value):  
    return tf.train.Feature(int64_list=tf.train.Int64List(value=[value]))  
  
# For representing float values. Useful for regression problems.  
def _float_feature(value):  
    return tf.train.Feature(float_list=tf.train.FloatList(value=[value]))
```

Write data as a TFRecord

```
# 1. Open a TFRecord file for writing  
with tf.io.TFRecordWriter(filename) as writer:  
    ...
```

Write data as a TFRecord

```
# 2. Create a features dictionary entry with the column names (features)
# as keys and the data wrapped with the types defined before.
#
# Use your own logic and python utilities to encode as string or
# bytes the structured data before passing it to the feature wrappers.
features=tf.train.Features(
    feature={
        "height": _int64_feature(image.shape[0]),
        "width": _int64_feature(image.shape[1]),
        "image": _bytes_feature(image.tostring()),
        "label": _bytes_feature(label.encode(encoding="utf-8"))
    }
)
```

Write data as a TFRecord

```
# 3. Wrap the feature dictionary within an `Example` which will be  
# a record on the TFRecord file.  
# Here `example` is a protocol buffer message.  
example = tf.train.Example(features)
```

¹[Protocol buffers](#) are a mechanism for serializing structured data developed by Google

Write data as a TFRecord

```
# 4. Serialize and write the example on the TFRecord file.  
writer.write(example.SerializeToString())
```

Write data as a TFRecord

```
with tf.io.TFRecordWriter(filename) as writer:
    for image, label in data:
        example = tf.train.Example(
            features=tf.train.Features(
                feature={
                    "height": _int64_feature(image.shape[0]),
                    "width": _int64_feature(image.shape[1]),
                    "image": _bytes_feature(image.tostring()),
                    "label": _bytes_feature(label.encode(encoding="utf-8"))
                }
            )
        )
        writer.write(example.SerializeToString())
```

Read data from a TFRecord file

```
# Data streamed from a TFRecord needs to be decoded back to numerical  
# types. This is done as part of the input pipeline with the help of  
# the `map` method of the `Dataset` objects.  
dataset = tf.data.TFRecordDataset(filename)  
dataset = dataset.map(decode)  
dataset = ...
```

Read data from a TFRecord file

```
# 1. Define a parser with the features you wish to extract and
# specify their types. Pass the serialized example to the parser.
example = tf.parse_single_example(
    serialized_example,
    features={
        "image": tf.FixedLenFeature([], tf.string),
        "label": tf.FixedLenFeature([], tf.string),
    })
```

Read data from a TFRecord file

```
# 2. Cast/decode each feature to the proper types.  
label = tf.cast(example["label"], tf.string)  
image = tf.decode_raw(example["image"], tf.uint8)
```

Read data from a TFRecord file

```
def decode(serialized_example):  
    features = tf.parse_single_example(  
        serialized_example,  
        features={  
            "height": tf.FixedLenFeature([], tf.int64),  
            "width": tf.FixedLenFeature([], tf.int64),  
            "image": tf.FixedLenFeature([], tf.string),  
            "label": tf.FixedLenFeature([], tf.string),  
        })  
    label = tf.cast(features["label"], tf.string)  
    width = tf.cast(features["width"], tf.int64)  
    height = tf.cast(features["height"], tf.int64)  
    image = tf.decode_raw(features["image"], tf.uint8)  
    image = tf.reshape(image, (height, width, 3))  
    return image, label
```

[lab] Reading and writing TFRecord files

Let's run the notebook `read_and_write_TFRecord_files.ipynb`. We are going to write two cat images to a TFRecord file and then we are going to read them to check that they can be recovered correctly.

Following the same steps, write the MNIST dataset to a TFRecord file. We will use it later on a lab.

You can get the MNIST data as numpy arrays with

```
from tensorflow.keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Best practices for storing data

- Avoid storing data as multiple small files.
- Divide large datasets ($\gg 1$ Gb) into TFRecord files of about 150 Mb.
- For smaller data sets (200MB-1GB), a single TFRecord is enough.
- Randomly divide the data into multiple files and then shuffle the filenames uniformly. This helps to decrease the size of the shuffle buffer.

[lab] Decoding an ImageNet TFRecord file

ImageNet is a large visual database designed for use in visual object recognition software research. It contains more than 14 million labeled images.

- Write an input pipeline reading one of the TFRecord files of the ImageNet dataset. You can find them at `/scratch/snx3000/stud71/imagenet`

The identifiers for the images and the labels are `'image/class/label'` (string) and `'image/encoded'` (int64) respectively. These images are encoded from jpeg format. You can use the functions `tf.image.decode_jpeg()` and `tf.image.resize_images()` of the `tf.image` API to decode them and resize them to `(224,224)`.

- Visualize some images and check that they were recovered correctly.
- Loop over the dataset and find how many images the TFRecord that you chose has, and how many different classes

Finding bottlenecks on the input pipeline

- Monitor GPU utilization with `nvidia-smi -l2`. A constant GPU utilization of 80-100% is expected for large models.
- Compare the number of samples per second given by your model with the one given by a very simple one, using the same input pipeline.
- Generate a timeline and check if there are blank spaces

Optimizing pipelining: Parallelize file reading

```
dataset = tf.data.Dataset.TFRecordDataset(['train1.tfrecords',  
                                            'train2.tfrecords',  
                                            ...],  
                                            num_parallel_reads=12)
```

Optimizing pipelining: Parallelize file reading

```
dataset = tf.data.Dataset.list_files(['train1.tfrecords',  
                                     'train2.tfrecords', ...])  
  
dataset = dataset.interleave(tf.data.TFRecordDataset,  
                             cycle_length=120,  
                             block_length=1,  
                             num_parallel_calls=12)
```

Optimizing pipelining: Parallelize mapping

```
# write vectorized map functions and map after batching  
# to take dvantage of vectorial operations  
dataset = dataset.map(decode,  
                      num_parallel_calls=12)
```

Optimizing pipelining: Parallelize mapping

```
# set the number of parallel calls dynamically at run time
dataset = dataset.map(decode,
                      num_parallel_calls=tf.data.experimental.AUTOTUNE)
```

Optimizing pipelining: Consider using fused transformations

```
# functionally, this is equivalent to a map followed by batch but
# may be more efficient than the two operations done individually
# tf.data.experimental.map_and_batch(
#     function,
#     batch_size,
#     num_parallel_batches=None,
#     drop_remainder=False,
#     num_parallel_calls=None)
```

```
dataset = dataset.apply(tf.data.experimental.map_and_batch(...))
```

Optimizing pipelining: Consider using fused transformations

```
# functionally, it is equivalent to a shuffle followed by repeat but
# may be more efficient than the two operations done individually
# tf.data.experimental.shuffle_and_repeat(
#     buffer_size,
#     count=None,
#     seed=None
# )

dataset = dataset.apply(tf.data.experimental.shuffle_and_repeat(...))
```


Optimizing pipelining: Consider `parallel_interleave`

```
# if a deterministic sequence on the iteration over the dataset
# is not necessary, sloppy interleave can enable additional
# performance optimizations
tf.data.experimental.parallel_interleave(
    map_func,
    cycle_length,
    block_length=1,
    sloppy=False,
    buffer_output_elements=None,
    prefetch_input_elements=None
)

# dataset = dataset.apply(tf.data.experimental.parallel_interleave(...))
```

Optimizing pipelining: Prefetching elements of the dataset

```
# prefetches elements of the dataset to the host memory.  
#  
# this can be set manually or dynamically at run time  
# with tf.data.experimental.AUTOTUNE  
dataset = dataset.prefetch(4)
```

Optimizing pipelining: Prefetching elements of the dataset

```
# prefetches elements of the dataset to the memory of the given device.  
# if used, this must be the final transformation in the input pipeline.  
dataset = dataset.apply(  
    tf.data.experimental.prefetch_to_device('/gpu:0',  
                                           buffer_size=4))
```

Feeding models with input pipelines (custom model)

```
# define the input pipeline
dataset = tf.data.Dataset.from_tensor_slices((x_numpy.astype(np.float32),
                                              y_numpy.astype(np.float32)))

dataset = dataset.batch(BATCH_SIZE)
dataset = dataset.repeat(NUM_EPOCHS)
iterator = dataset.make_one_shot_iterator()
next_item = iterator.get_next()
```

Feeding models with input pipelines (custom model)

```
# define the model ( $y = m * x + n$ )
slope = tf.Variable(np.random.randn() , name='slope')
offset = tf.Variable(np.random.randn() , name='offset')
x, y = next_item
y_hat = slope * x + offset

# define the loss
loss = tf.losses.mean_squared_error(y_hat, y)

# choose an optimizer and create minimization and
# variable initialization operations
...
```

Feeding models with input pipelines (custom model)

```
# define the input pipeline
```

```
dataset = tf.data.Dataset.from_tensor_slices((x_numpy.astype(np.float32),  
                                              y_numpy.astype(np.float32)))
```

```
dataset = dataset.batch(BATCH_SIZE)
```

```
dataset = dataset.repeat(NUM_EPOCHS)
```

```
iterator = dataset.make_one_shot_iterator()
```

```
x, y = iterator.get_next()
```

Feeding models with input pipelines (custom model)

```
# define the model ( $y = m * x + n$ )
slope = tf.Variable(np.random.randn() , name='slope')
offset = tf.Variable(np.random.randn() , name='offset')
y_hat = slope * x + offset

# define the loss
loss = tf.losses.mean_squared_error(y_hat, y)

# choose an optimizer and create minimization and
# variable initialization operations
...
```

[lab] Feeding models with `tf.data` input pipelines

Let's run the notebooks in `input_pipelines/feeding_models` and check out how to feed models using Keras and the `tf.estimator` API.

TensorFlow Datasets

TensorFlow Datasets provides many public datasets as `tf.data.Datasets` and you can contribute back with your datasets.

<https://github.com/tensorflow/datasets>

[video] TensorFlow Datasets (TensorFlow Dev Summit 2019)

Useful info on input pipelines

- [Importing data into TensorFlow](#)
- [Input Pipelines Performance](#)
- [General best practices for performance in TensorFlow](#)
- [\[video\] tf.data: Fast, flexible, and easy-to-use input pipelines](#)
(TensorFlow Dev Summit 2018)
- [\[video\] Training Performance: A user's guide to converge faster](#)
(TensorFlow Dev Summit 2018)

Thank you for your attention!